

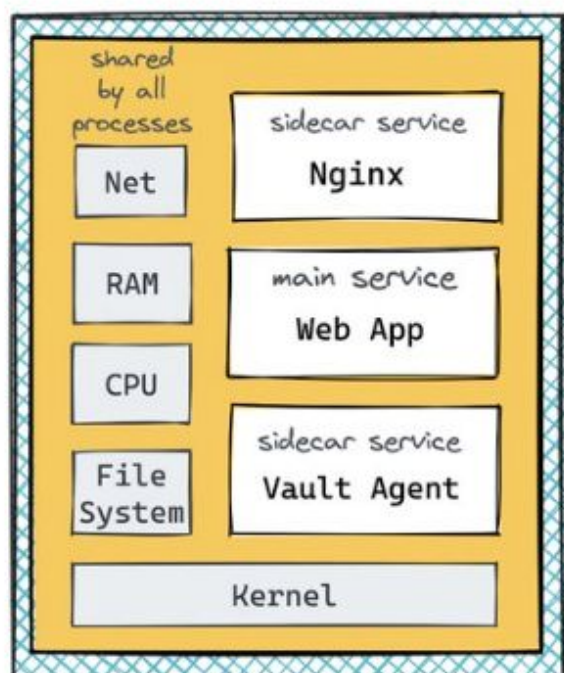
Kubernetes basics explained by analogy 🧵

...or "How Kubernetes Just Repeats Good Old Deployment Patterns"

1. For a long time, people had been deploying services as groups of virtual (or physical) machines.

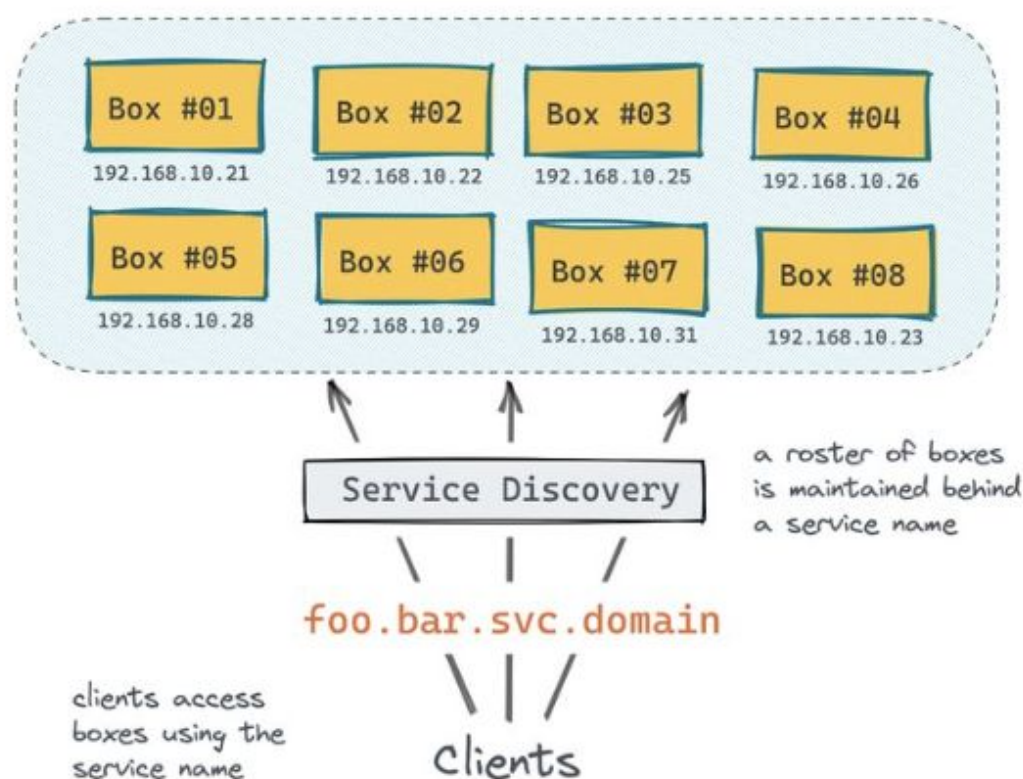
But VMs were often slow and bulky. Hence, not very efficient.

Virtual Machine - a "Box" ...or real!



typically has a dedicated address
e.g. 192.168.10.5

Service - a named group of identical "Boxes" distributed!

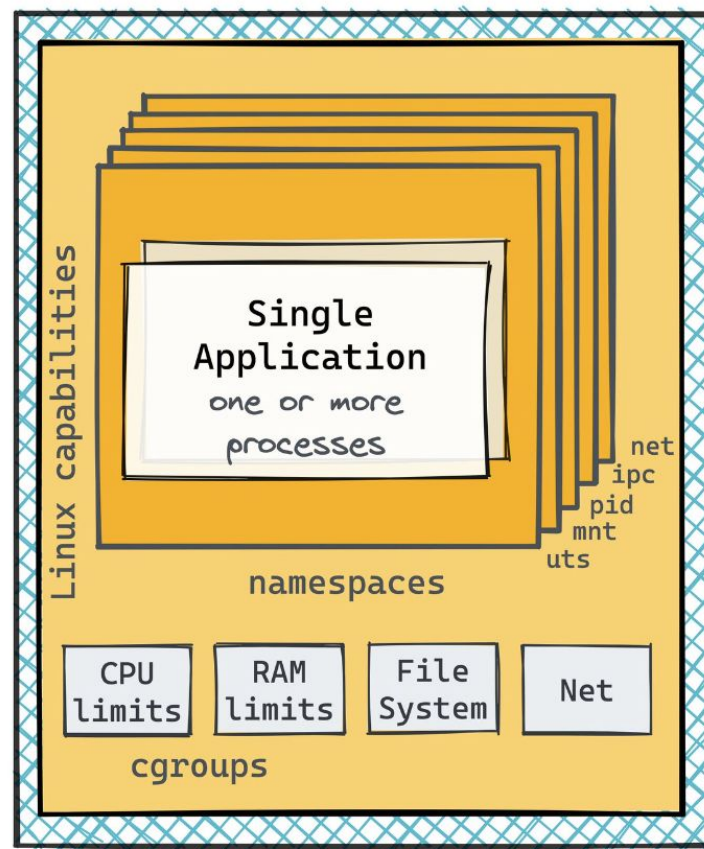


2. Then containers gained quite some popularity.

With containers, it became easier to distribute services. Reproducibility also improved. But containers haven't become a replacement for VMs.

Mainly, because of their deliberate focus on being an environment to run a single app:

OCI Container - a new "Box"?



typically has a dedicated address

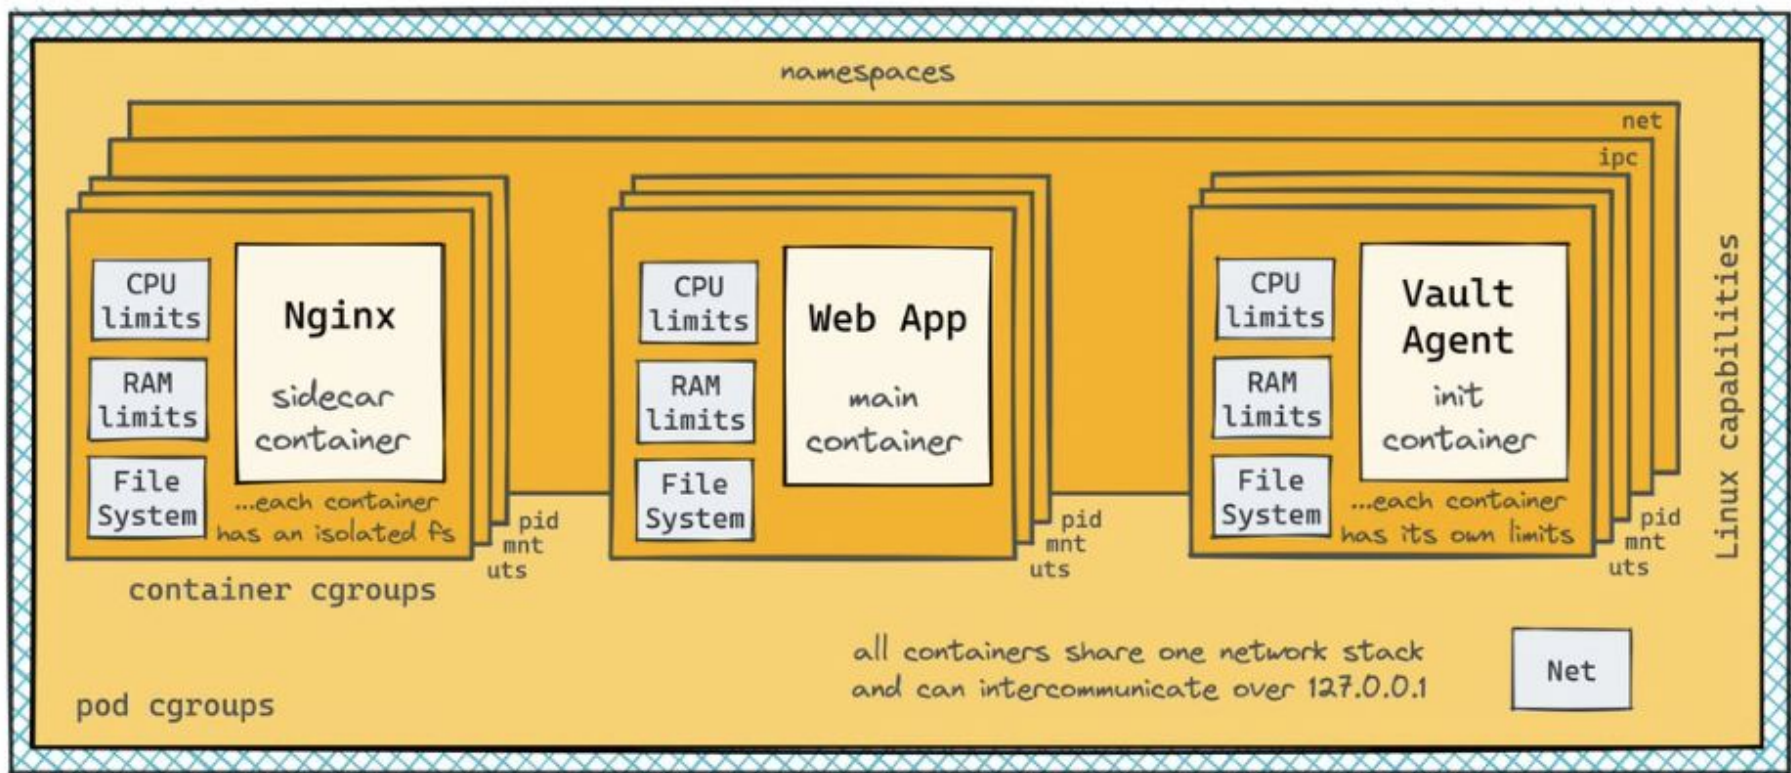
e.g. 172.18.0.3

3. Instead of containers, another abstraction took off - Kubernetes Pods!

A Pod is a group of semi-fused containers. External borders were preserved, but some of the internal isolation b/w containers substituting a Pod got weakened.

A Pod is a much closer abstraction to a VM.

Kubernetes Pod - the new "Box"!



Pod - a group of "semi-fused" containers

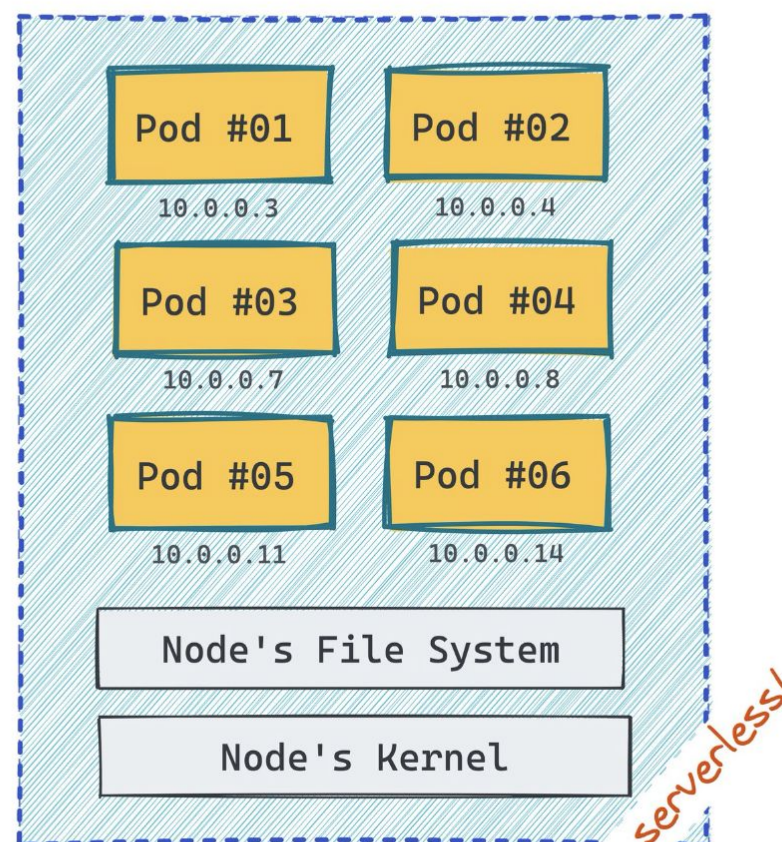
all containers in a Pod are addressable via a single IP address. E.g. 10.0.0.3

4. A single (virtual or physical) machine can run many independant Pods.

In Kubernetes, machines substituting a cluster are called Nodes, but developers are rarely concerned with this abstraction. For them, Kubernetes is serverless!

More Pods per server means better packing.

Kubernetes Node - "invisible" for Developers



Nodes and Pods often live in different subnets

192.168.10.5

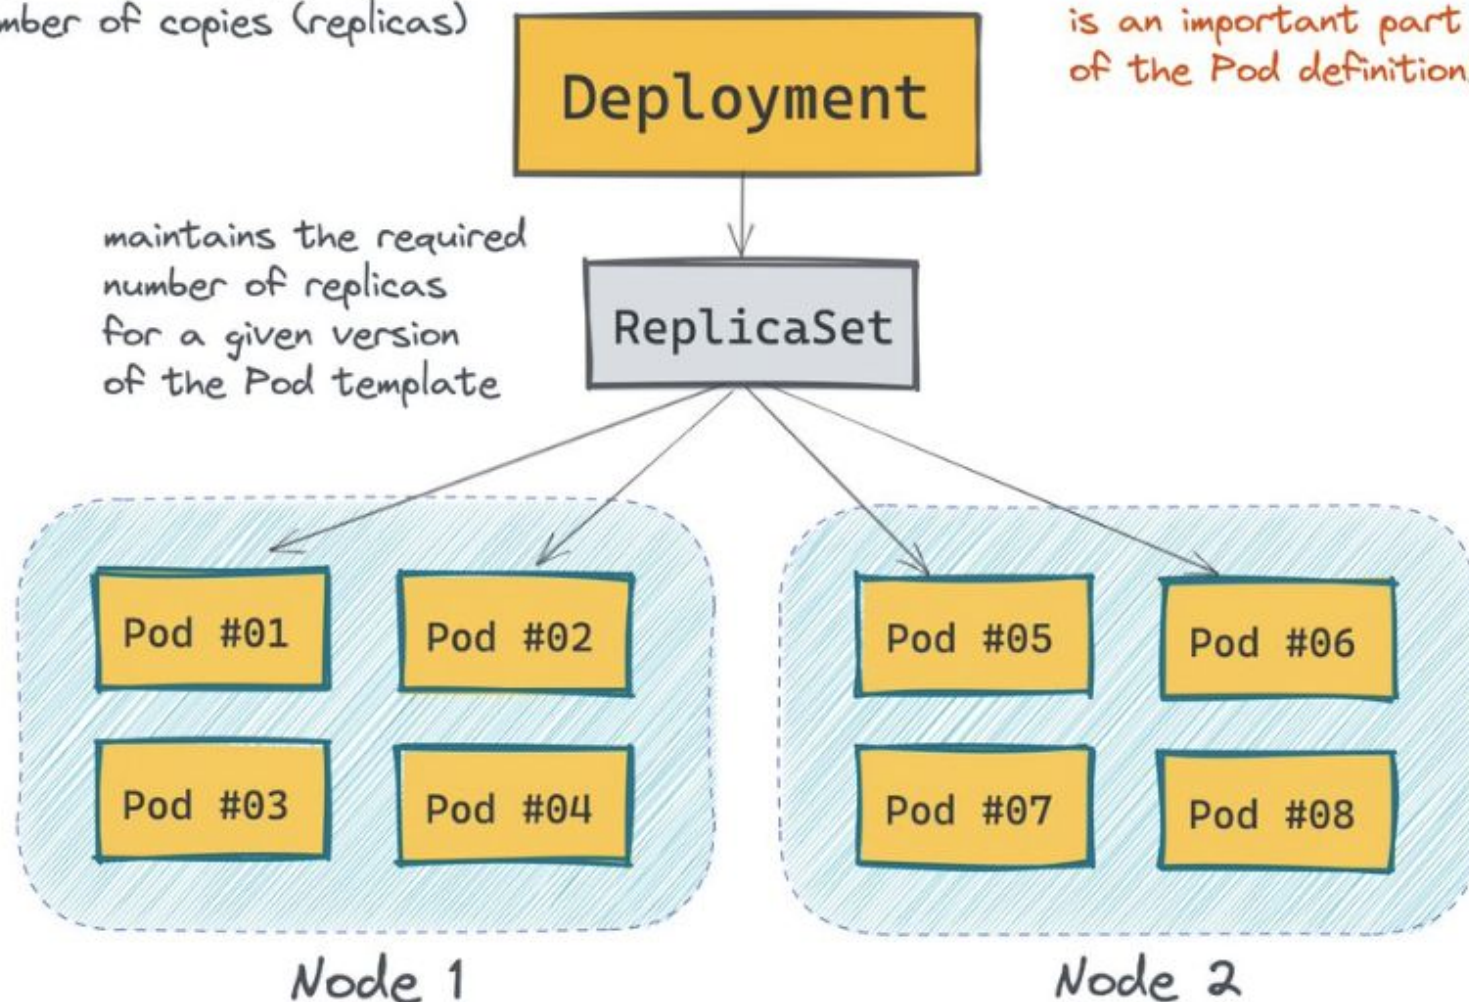
5. Deployment of Pods happens through replicating a Pod template.

There is a Deployment object in Kubernetes that holds the desired Pod template and the needed number of "copies." But logically, there is not much difference between scaling Pods and VMs.

Kubernetes Deployment - a means to replicate "Boxes"

defines a Pod template and
a number of copies (replicas)

`Pod.metadata.labels`
is an important part
of the Pod definition!




6. Kubernetes Service is a means of grouping Pods behind a logical name.

Kubernetes comes with built-in service discovery.

The implementation is neither client- nor server-side (rather network-side). But from the clients' standpoint, it feels like a good old reverse proxy.

Kubernetes Built-In Service Discovery



Kubernetes Service - a logical group of similarly-labeled Pods

